

Explanation of the MapReduce process, followed by a flowchart description.

MapReduce Overview

1. Map Phase:

- Input data is divided into splits, and the mapper processes each split to produce intermediate key-value pairs.

2. Partitioning:

- Intermediate key-value pairs are assigned to different reducers based on a partitioning function (typically a hash function).

3. Shuffling:

- The shuffle phase sorts and transfers the intermediate data from mappers to the appropriate reducers. This step involves sorting the keys.

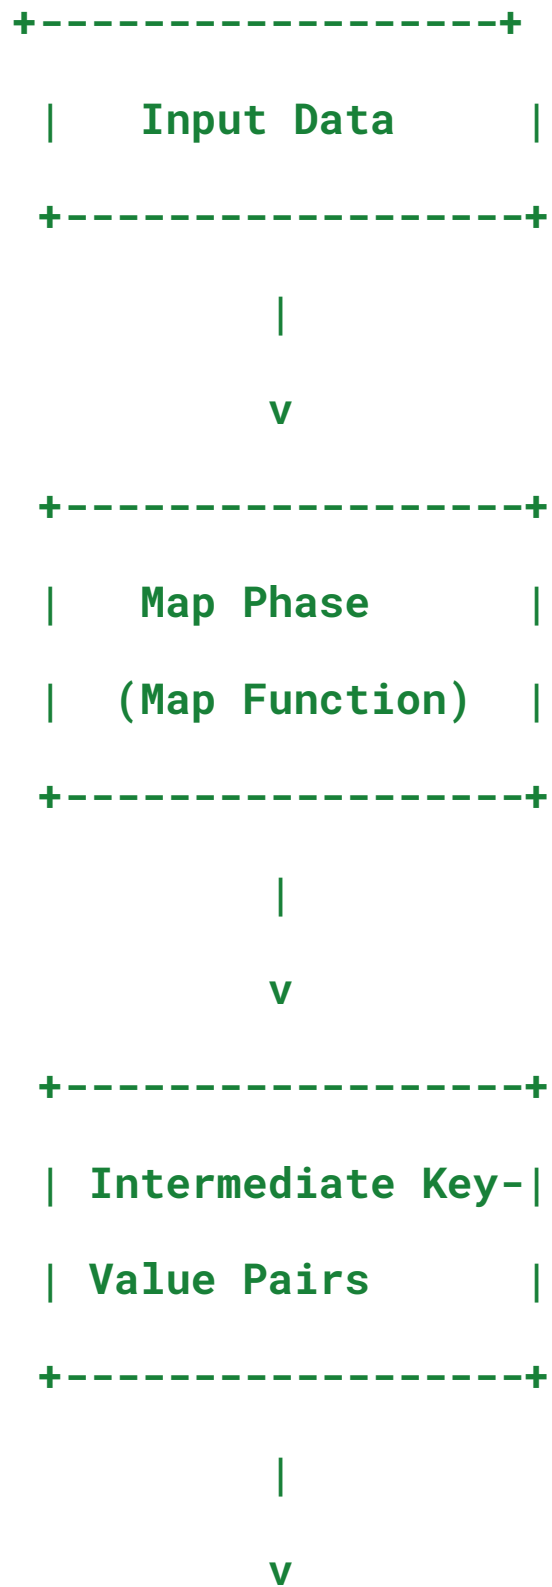
4. Reduce Phase:

- Reducers aggregate the values associated with each key to produce final output.

Flowchart for Partitioning and Shuffling in MapReduce

Here's a textual representation of the flowchart:

Copy code



+-----+

| Partitioning |

| (Assigning to |

| Reducers) |

+-----+

|

v

+-----+

| Shuffling |

| (Sorting and |

| Transferring) |

+-----+

|


v

+-----+

| Reduce Phase |

| (Reduce Function) |

+-----+



Explanation of Each Step

1. Input Data:

- The raw input data is fed into the MapReduce framework.

2. Map Phase:

- The data is processed in the map phase where each mapper generates intermediate key-value pairs.

3. Intermediate Key-Value Pairs:

- The output of the map phase, consisting of many key-value pairs.

4. Partitioning:

- The intermediate key-value pairs are assigned to reducers based on the partitioning logic (e.g., hashing keys).

5. Shuffling:

- The intermediate data is sorted by key and transferred to the respective reducers. This ensures that all values for the same key are sent to the same reducer.

6. Reduce Phase:

- Each reducer processes the received key-value pairs to produce a smaller set of results.

7. Final Output:

- The results from the reduce phase are combined to produce the final output of the MapReduce job.

How Map reduce managed by YARN

In YARN (Yet Another Resource Negotiator), which is a key component of the Hadoop ecosystem, the execution of MapReduce jobs—including the Map, Partitioning, Shuffling, and Reduce phases—occurs in a more structured and resource-efficient manner compared to earlier Hadoop implementations. Here's how these phases are executed in YARN:

Overview of YARN Architecture

YARN separates the resource management and job scheduling from the processing framework. It consists of three main components:

1. **ResourceManager (RM)**: Manages the cluster resources and schedules jobs.
2. **NodeManager (NM)**: Manages the resources on a single node and handles the execution of tasks.
3. **ApplicationMaster (AM)**: Manages the lifecycle of a single application, negotiating resources with the ResourceManager and working with the NodeManagers to execute tasks.

Execution Phases in YARN

1. Job Submission

- The user submits a MapReduce job through the **YARN Command Line Interface** or through an API.
- The submission includes the application details, such as the jar file, input data location, and output data location.

2. Resource Allocation

- The ResourceManager receives the job submission request.
- It allocates resources for the ApplicationMaster (AM), which will manage the execution of the MapReduce job.
- The ResourceManager selects a NodeManager to launch the ApplicationMaster based on available resources.

3. ApplicationMaster Starts

- The NodeManager starts the ApplicationMaster in a separate container.
- The ApplicationMaster registers itself with the ResourceManager and requests resources for the map and reduce tasks.

4. Map Phase Execution

- The ApplicationMaster requests resources for the mappers from the ResourceManager.

- The ResourceManager allocates containers for the mappers on suitable NodeManagers.
- The NodeManagers launch the mapper tasks in their containers.
- Each mapper processes input data splits, generating intermediate key-value pairs.

5. Partitioning

- After mappers complete their tasks, the intermediate data is partitioned based on the partitioning function.
- Each mapper sends its intermediate key-value pairs to the corresponding reducers, using the defined partitioning logic.

6. Shuffling Phase

- During the shuffling phase, the intermediate key-value pairs are sorted by key and transferred from the mappers to the reducers.
- This process is managed by the ApplicationMaster, which tracks data transfer and ensures that all data for a particular key is sent to the same reducer.
- The shuffling process can involve **disk I/O**, where mappers write intermediate data to disk before transferring it to reducers.

7. Reduce Phase Execution

- The ApplicationMaster requests resources for the reducers in a similar way to the mappers.
- Once the ResourceManager allocates resources, the NodeManagers launch the reducer tasks in their containers.
- Each reducer processes the received key-value pairs, aggregating the values associated with each key.

8. Final Output Generation

- After the reducers complete their tasks, they write the final output to the specified output location.
- The ApplicationMaster monitors the completion of all tasks and handles any failures, restarting tasks if necessary.
- Once all tasks are finished, the ApplicationMaster informs the ResourceManager that the job is complete.

9. Job Cleanup

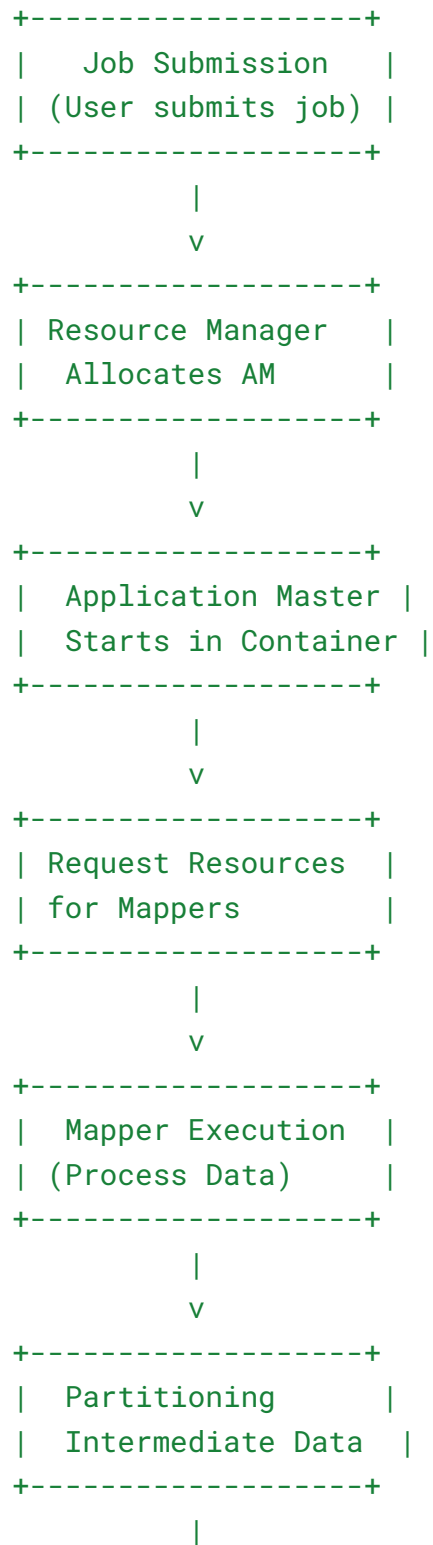
- The ApplicationMaster handles cleanup operations, releasing resources allocated for the job.
- The NodeManagers also clean up the containers used for the tasks.

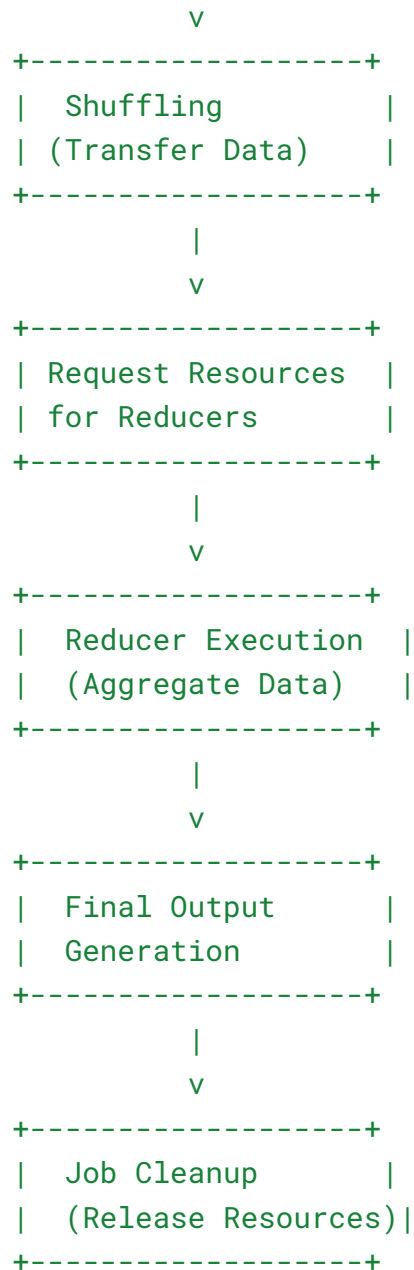
Flowchart of YARN Execution Phases

Here's a simplified flowchart representing the execution of MapReduce phases in YARN:

sql

Copy code





Summary

In summary, YARN effectively manages the execution of MapReduce jobs by coordinating between ResourceManager, NodeManagers, and ApplicationMaster. This separation allows for better resource utilization, scalability, and fault tolerance compared to earlier Hadoop versions. The overall architecture and execution flow ensure that each phase of MapReduce is executed efficiently and in a distributed manner.

